Artificial Intelligence Nanodegree

Convolutional Neural Networks

Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with **'(IMPLEMENTATION)'** in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

> **Note**: Once you have completed all of the code implementations, you need to finalize your work by exporting the iPython Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to \n", "**File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a **'Question X'** header. Carefully read each question and provide thorough answers in the following text boxes that begin with **'Answer:'**. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

> **Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this IPython notebook.

Why We're Here

In this notebook, you will make the first steps towards developing an algorithm that could be used as part of a mobile or web app. At the end of this project, your code will accept any user-supplied image as input. If a dog is detected in the image, it will provide an estimate of the dog's breed. If a human is detected, it will provide an estimate of the dog breed that is most resembling. The image below displays potential sample output of your finished project (... but we expect that each student's algorithm will behave differently!).

```
hello, dog!
your predicted breed is ...
American Staffordshire terrier
```


In this real-world setting, you will need to piece together a series of models to perform different tasks; for instance, the algorithm that detects humans in an image will be different from the CNN that infers dog breed. There are many points of possible failure, and no perfect algorithm exists. Your imperfect solution will nonetheless create a fun user experience!

The Road Ahead

We break the notebook into separate steps. Feel free to use the links below to navigate the notebook.

- Step 0: Import Datasets
- Step 1: Detect Humans
- Step 2: Detect Dogs
- Step 3: Create a CNN to Classify Dog Breeds (from Scratch)
- Step 4: Use a CNN to Classify Dog Breeds (using Transfer Learning)
- Step 5: Create a CNN to Classify Dog Breeds (using Transfer Learning)
- Step 6: Write your Algorithm
- Step 7: Test Your Algorithm

Step 0: Import Datasets

Import Dog Dataset

In the code cell below, we import a dataset of dog images. We populate a few variables through the use of the `load_files` function from the scikit-learn library:

- `train_files`, `valid_files`, `test_files` - numpy arrays containing file paths to images
- `train_targets`, `valid_targets`, `test_targets` - numpy arrays containing onehot-encoded classification labels
- `dog_names` - list of string-valued dog breed names for translating labels

```
In [1]:  from sklearn.datasets import load_files
         from keras.utils import np_utils
         import numpy as np
         from glob import glob

         # define function to load train, test, and validation datasets
         def load_dataset(path):
             data = load_files(path)
             dog_files = np.array(data['filenames'])
             dog_targets = np_utils.to_categorical(np.array(data['target']), 133)
             return dog_files, dog_targets

         # load train, test, and validation datasets
         train_files, train_targets = load_dataset('dogImages/train')
         valid_files, valid_targets = load_dataset('dogImages/valid')
         test_files, test_targets = load_dataset('dogImages/test')

         # load list of dog names
         dog_names = [item[20:-1] for item in sorted(glob("dogImages/train/*/"))]

         # print statistics about the dataset
         print('There are %d total dog categories.' % len(dog_names))
         print('There are %s total dog images.\n' % len(np.hstack([train_files, v
         alid_files, test_files])))
         print('There are %d training dog images.' % len(train_files))
         print('There are %d validation dog images.' % len(valid_files))
         print('There are %d test dog images.'% len(test_files))
```

```
Using TensorFlow backend.

There are 133 total dog categories.
There are 8351 total dog images.

There are 6680 training dog images.
There are 835 validation dog images.
There are 836 test dog images.
```

Import Human Dataset

In the code cell below, we import a dataset of human images, where the file paths are stored in the numpy array `human_files`.

```
In [2]:  import random
         random.seed(8675309)

         # load filenames in shuffled human dataset
         human_files = np.array(glob("lfw/*/*"))
         random.shuffle(human_files)

         # print statistics about the dataset
         print('There are %d total human images.' % len(human_files))
```

```
There are 13233 total human images.
```

Step 1: Detect Humans

We use OpenCV's implementation of Haar feature-based cascade classifiers (http://docs.opencv.org/trunk/d7/d8b/tutorial_py_face_detection.html) to detect human faces in images. OpenCV provides many pre-trained face detectors, stored as XML files on github (https://github.com/opencv/opencv/tree/master/data/haarcascades). We have downloaded one of these detectors and stored it in the `haarcascades` directory.

In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [3]:  import cv2
         import matplotlib.pyplot as plt
         %matplotlib inline

         # extract pre-trained face detector
         face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalfa
         ce_alt.xml')

         # load color (BGR) image
         img = cv2.imread(human_files[3])
         # convert BGR image to grayscale
         gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

         # find faces in image
         faces = face_cascade.detectMultiScale(gray)

         # print number of faces detected in the image
         print('Number of faces detected:', len(faces))

         # get bounding box for each detected face
         for (x,y,w,h) in faces:
             # add bounding box to color image
             cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

         # convert BGR image to RGB for plotting
         cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

         # display the image, along with bounding box
         plt.imshow(cv_rgb)
         plt.show()
```

Number of faces detected: 3

Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [4]:   # returns "True" if face is detected in image stored at img_path
          def face_detector(img_path):
              img = cv2.imread(img_path)
              gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
              faces = face_cascade.detectMultiScale(gray)
              return len(faces) > 0
```

(IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer:

The `face_detector` function detected a face 99% of the images in `human_files` and 11% of the images in the `dog_files`. It is clearly better at detecting human faces, that it is at not detecting human faces.

```
In [5]: human_files_short = human_files[:100]
        dog_files_short = train_files[:100]
        # Do NOT modify the code above this line.

        ## TODO: Test the performance of the face_detector algorithm
        ## on the images in human_files_short and dog_files_short.

        h = sum([face_detector(h) for h in human_files_short]) / 100.
        d = sum([face_detector(d) for d in dog_files_short]) / 100.

        print('face_dectector detected human faces in %f of human images' % h)
        print('face_dectector detected human faces in %f of dog images' % d)
```

```
face_dectector detected human faces in 0.990000 of human images
face_dectector detected human faces in 0.110000 of dog images
```

Question 2: This algorithmic choice necessitates that we communicate to the user that we accept human images only when they provide a clear view of a face (otherwise, we risk having unneccessarily frustrated users!). In your opinion, is this a reasonable expectation to pose on the user? If not, can you think of a way to detect humans in images that does not necessitate an image with a clearly presented face?

Answer:

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on each of the datasets.

```
In [6]: ## (Optional) TODO: Report the performance of another
        ## face detection algorithm on the LFW dataset
        ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a pre-trained ResNet-50 (http://ethereon.github.io/netscope/#/gist/db945b393d40bfa26006) model to detect dogs in images. Our first line of code downloads the ResNet-50 model, along with weights that have been trained on ImageNet (http://www.image-net.org/), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories (https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a). Given an image, this pre-trained ResNet-50 model returns a prediction (derived from the available categories in ImageNet) for the object that is contained in the image.

```
In [7]: from keras.applications.resnet50 import ResNet50

        # define ResNet50 model
        ResNet50_model = ResNet50(weights='imagenet')
```

Pre-process the Data

When using TensorFlow as backend, Keras CNNs require a 4D array (which we'll also refer to as a 4D tensor) as input, with shape

$$(nb\_samples, rows, columns, channels),$$

where `nb_samples` corresponds to the total number of images (or samples), and `rows`, `columns`, and `channels` correspond to the number of rows, columns, and channels for each image, respectively.

The `path_to_tensor` function below takes a string-valued file path to a color image as input and returns a 4D tensor suitable for supplying to a Keras CNN. The function first loads the image and resizes it to a square image that is 224×224 pixels. Next, the image is converted to an array, which is then resized to a 4D tensor. In this case, since we are working with color images, each image has three channels. Likewise, since we are processing a single image (or sample), the returned tensor will always have shape

$$(1, 224, 224, 3).$$

The `paths_to_tensor` function takes a numpy array of string-valued image paths as input and returns a 4D tensor with shape

$$(nb\_samples, 224, 224, 3).$$

Here, `nb_samples` is the number of samples, or number of images, in the supplied array of image paths. It is best to think of `nb_samples` as the number of 3D tensors (where each 3D tensor corresponds to a different image) in your dataset!

```
In [8]:  from keras.preprocessing import image
         from tqdm import tqdm

         def path_to_tensor(img_path):
             # loads RGB image as PIL.Image.Image type
             img = image.load_img(img_path, target_size=(224, 224))
             # convert PIL.Image.Image type to 3D tensor with shape (224, 224, 3)
             x = image.img_to_array(img)
             # convert 3D tensor to 4D tensor with shape (1, 224, 224, 3) and ret
         urn 4D tensor
             return np.expand_dims(x, axis=0)

         def paths_to_tensor(img_paths):
             list_of_tensors = [path_to_tensor(img_path) for img_path in tqdm(img
         _paths)]
             return np.vstack(list_of_tensors)
```

Making Predictions with ResNet-50

Getting the 4D tensor ready for ResNet-50, and for any other pre-trained model in Keras, requires some additional processing. First, the RGB image is converted to BGR by reordering the channels. All pre-trained models have the additional normalization step that the mean pixel (expressed in RGB as [103.939, 116.779, 123.68] and calculated from all pixels in all images in ImageNet) must be subtracted from every pixel in each image. This is implemented in the imported function `preprocess_input`. If you're curious, you can check the code for `preprocess_input` here (https://github.com/fchollet/keras/blob/master/keras/applications/imagenet_utils.py).

Now that we have a way to format our image for supplying to ResNet-50, we are now ready to use the model to extract the predictions. This is accomplished with the `predict` method, which returns an array whose i-th entry is the model's predicted probability that the image belongs to the i-th ImageNet category. This is implemented in the `ResNet50_predict_labels` function below.

By taking the argmax of the predicted probability vector, we obtain an integer corresponding to the model's predicted object class, which we can identify with an object category through the use of this dictionary (https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a).

```
In [9]: from keras.applications.resnet50 import preprocess_input, decode_predict
        ions

        def ResNet50_predict_labels(img_path):
            # returns prediction vector for image located at img_path
            img = preprocess_input(path_to_tensor(img_path))
            return np.argmax(ResNet50_model.predict(img))
```

Write a Dog Detector

While looking at the dictionary (https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from `'Chihuahua'` to `'Mexican hairless'`. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained ResNet-50 model, we need only check if the `ResNet50_predict_labels` function above returns a value between 151 and 268 (inclusive).

We use these ideas to complete the `dog_detector` function below, which returns `True` if a dog is detected in an image (and `False` if not).

```
In [10]: ### returns "True" if a dog is detected in the image stored at img_path
         def dog_detector(img_path):
             prediction = ResNet50_predict_labels(img_path)
             return ((prediction <= 268) & (prediction >= 151))
```

(IMPLEMENTATION) Assess the Dog Detector

Question 3: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer:

The `dog_detector` detected a dog in 0% of the human images, and it detected 100% of the dogs in the dog images. It is a great detector of dogs in images.

```
In [12]:   ### TODO: Test the performance of the dog_detector function
           ### on the images in human_files_short and dog_files_short.
           h = sum([dog_detector(h) for h in human_files_short]) / 100.
           d = sum([dog_detector(d) for d in dog_files_short]) / 100.

           print('dog_detector detected a dog in %f of human images' % h)
           print('dog_detector detected a dog in %f of dog images' % d)
```

```
dog_detector detected a dog in 0.000000 of human images
dog_detector detected a dog in 1.000000 of dog images
```

# Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet*!), and you must attain a test accuracy of at least 1%. In Step 5 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

Be careful with adding too many trainable layers! More parameters means longer training, which means you are more likely to need a GPU to accelerate the training process. Thankfully, Keras provides a handy estimate of the time that each epoch is likely to take; you can extrapolate this estimate to figure out how long it will take for your algorithm to train.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have great difficulty in distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
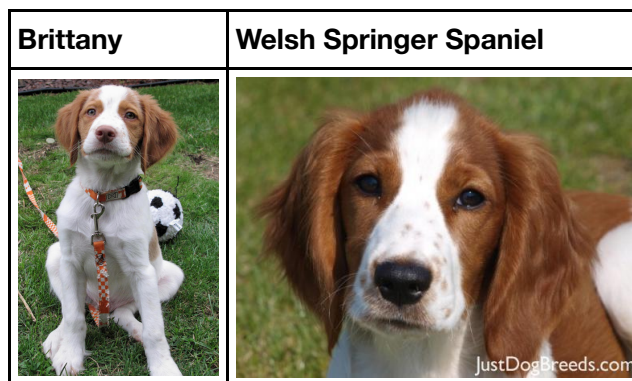
It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
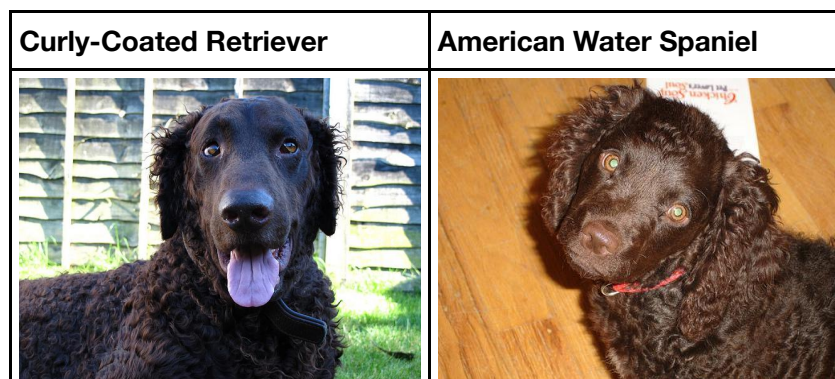
Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador	Black Labrador

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imabalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

Pre-process the Data

We rescale the images by dividing every pixel in every image by 255.

```
In [10]:  from PIL import ImageFile
          ImageFile.LOAD_TRUNCATED_IMAGES = True

          # pre-process the data for Keras
          train_tensors = paths_to_tensor(train_files).astype('float32')/255
          valid_tensors = paths_to_tensor(valid_files).astype('float32')/255
          test_tensors = paths_to_tensor(test_files).astype('float32')/255
```

```
100%|██████████| 6680/6680 [00:53<00:00, 124.93it/s]
100%|██████████| 835/835 [00:06<00:00, 139.01it/s]
100%|██████████| 836/836 [00:05<00:00, 139.78it/s]
```

(IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
model.summary()
```

We have imported some Python modules to get you started, but feel free to import as many modules as you need. If you end up getting stuck, here's a hint that specifies a model that trains relatively fast on CPU and attains >1% test accuracy in 5 epochs:

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 223, 223, 16)	208
max_pooling2d_1 (MaxPooling2	(None, 111, 111, 16)	0
conv2d_2 (Conv2D)	(None, 110, 110, 32)	2080
max_pooling2d_2 (MaxPooling2	(None, 55, 55, 32)	0
conv2d_3 (Conv2D)	(None, 54, 54, 64)	8256
max_pooling2d_3 (MaxPooling2	(None, 27, 27, 64)	0
global_average_pooling2d_1 (	(None, 64)	0
dense_1 (Dense)	(None, 133)	8645

```
Total params: 19,189.0
Trainable params: 19,189.0
Non-trainable params: 0.0
```

INPUT
CONV
POOL
CONV
POOL
CONV
POOL
GAP
DENSE

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. If you chose to use the hinted architecture above, describe why you think that CNN architecture should work well for the image classification task.

Answer:

I decided to use the above architecture because it has the basic components of an image classification model, with a series of convolution, max pooling and dropout layers put together, with one (sometimes more) fully conected layer at the end.

This general patter of convolutions layers and ending with fully conected layers has been proven to be a good for image recognition tasks in the literature. The basic idea being that the convolution layers extract the features from the images, and the dense layers build higher level features form the convolution features.

```
In [11]:  from keras.layers import Conv2D, MaxPooling2D, GlobalAveragePooling2D
          from keras.layers import Dropout, Flatten, Dense
          from keras.models import Sequential

          model = Sequential()

          ### TODO: Define your architecture.
          model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224,
          3)))
          model.add(MaxPooling2D(pool_size=(2, 2)))
          model.add(Dropout(0.2))

          model.add(Conv2D(32, (3, 3), activation='relu'))
          model.add(MaxPooling2D(pool_size=(2, 2)))
          model.add(Dropout(0.2))

          model.add(Flatten())

          model.add(Dense(133, activation='softmax'))


          model.summary()
```

```
_____
Layer (type)                 Output Shape              Param #
=================================================================
conv2d_1 (Conv2D)            (None, 222, 222, 32)      896
_____
max_pooling2d_2 (MaxPooling2 (None, 111, 111, 32)      0
_____
dropout_1 (Dropout)          (None, 111, 111, 32)      0
_____
conv2d_2 (Conv2D)            (None, 109, 109, 32)      9248
_____
max_pooling2d_3 (MaxPooling2 (None, 54, 54, 32)        0
_____
dropout_2 (Dropout)          (None, 54, 54, 32)        0
_____
flatten_2 (Flatten)          (None, 93312)             0
_____
dense_1 (Dense)              (None, 133)               12410629
=================================================================
Total params: 12,420,773
Trainable params: 12,420,773
Non-trainable params: 0
_____
```

Compile the Model

```
In [12]:  model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metr
          ics=['accuracy'])
```

(IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to augment the training data (https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.

In [13]:
```python
from keras.callbacks import ModelCheckpoint

### TODO: specify the number of epochs that you would like to use to tra
in the model.

epochs = 5

### Do NOT modify the code below this line.

checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.from_
scratch.hdf5',
                               verbose=1, save_best_only=True)

model.fit(train_tensors, train_targets,
          validation_data=(valid_tensors, valid_targets),
          epochs=epochs, batch_size=20, callbacks=[checkpointer], verbos
e=1)
```

```
Train on 6680 samples, validate on 835 samples
Epoch 1/5
6660/6680 [=============================>.] - ETA: 0s - loss: 4.9659 - a
cc: 0.0086Epoch 00000: val_loss improved from inf to 4.87889, saving mo
del to saved_models/weights.best.from_scratch.hdf5
6680/6680 [==============================] - 33s - loss: 4.9656 - acc:
 0.0085 - val_loss: 4.8789 - val_acc: 0.0084
Epoch 2/5
6660/6680 [=============================>.] - ETA: 0s - loss: 4.7507 - a
cc: 0.0258Epoch 00001: val_loss improved from 4.87889 to 4.78057, savin
g model to saved_models/weights.best.from_scratch.hdf5
6680/6680 [==============================] - 31s - loss: 4.7510 - acc:
 0.0257 - val_loss: 4.7806 - val_acc: 0.0251
Epoch 3/5
6660/6680 [=============================>.] - ETA: 0s - loss: 3.2858 - a
cc: 0.2761Epoch 00002: val_loss did not improve
6680/6680 [==============================] - 31s - loss: 3.2868 - acc:
 0.2754 - val_loss: 5.7846 - val_acc: 0.0491
Epoch 4/5
6660/6680 [=============================>.] - ETA: 0s - loss: 1.0563 - a
cc: 0.7580Epoch 00003: val_loss did not improve
6680/6680 [==============================] - 31s - loss: 1.0554 - acc:
 0.7584 - val_loss: 7.8313 - val_acc: 0.0491
Epoch 5/5
6660/6680 [=============================>.] - ETA: 0s - loss: 0.2872 - a
cc: 0.9333Epoch 00004: val_loss did not improve
6680/6680 [==============================] - 31s - loss: 0.2867 - acc:
 0.9334 - val_loss: 10.2375 - val_acc: 0.0527
```

Out[13]: <keras.callbacks.History at 0x7fbb82f25c50>

Load the Model with the Best Validation Loss

In [16]:
```python
model.load_weights('saved_models/weights.best.from_scratch.hdf5')
```

Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 1%.

```
In [17]: # get index of predicted dog breed for each image in test set
         dog_breed_predictions = [np.argmax(model.predict(np.expand_dims(tensor,
         axis=0))) for tensor in test_tensors]

         # report test accuracy
         test_accuracy = 100*np.sum(np.array(dog_breed_predictions)==np.argmax(te
         st_targets, axis=1))/len(dog_breed_predictions)
         print('Test accuracy: %.4f%%' % test_accuracy)
```

```
Test accuracy: 4.7847%
```

Step 4: Use a CNN to Classify Dog Breeds

To reduce training time without sacrificing accuracy, we show you how to train a CNN using transfer learning. In the following step, you will get a chance to use transfer learning to train your own CNN.

Obtain Bottleneck Features

```
In [18]: bottleneck_features = np.load('bottleneck_features/DogVGG16Data.npz')
         train_VGG16 = bottleneck_features['train']
         valid_VGG16 = bottleneck_features['valid']
         test_VGG16 = bottleneck_features['test']
```

Model Architecture

The model uses the the pre-trained VGG-16 model as a fixed feature extractor, where the last convolutional output of VGG-16 is fed as input to our model. We only add a global average pooling layer and a fully connected layer, where the latter contains one node for each dog category and is equipped with a softmax.

In [19]:
```
VGG16_model = Sequential()
VGG16_model.add(GlobalAveragePooling2D(input_shape=train_VGG16.shape[1:]))

VGG16_model.add(Dense(133, activation='softmax'))

VGG16_model.summary()
```

```
_____
Layer (type)                 Output Shape              Param #
=================================================================
global_average_pooling2d_1 ( (None, 512)               0
_____
dense_2 (Dense)              (None, 133)               68229
=================================================================
Total params: 68,229
Trainable params: 68,229
Non-trainable params: 0
_____
```

Compile the Model

In [20]:
```
VGG16_model.compile(loss='categorical_crossentropy',
optimizer='rmsprop', metrics=['accuracy'])
```

Train the Model

```
In [21]: checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.VGG1
         6.hdf5',
                                         verbose=1, save_best_only=True)

         VGG16_model.fit(train_VGG16, train_targets,
                 validation_data=(valid_VGG16, valid_targets),
                 epochs=20, batch_size=20, callbacks=[checkpointer], verbose=1)
```

```
Train on 6680 samples, validate on 835 samples
Epoch 1/20
6600/6680 [============================>.] - ETA: 0s - loss: 12.1850 -
 acc: 0.1276Epoch 00000: val_loss improved from inf to 10.58423, saving
model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 1s - loss: 12.1651 - acc:
 0.1290 - val_loss: 10.5842 - val_acc: 0.2287
Epoch 2/20
6640/6680 [============================>.] - ETA: 0s - loss: 9.9561 - a
cc: 0.2956 Epoch 00001: val_loss improved from 10.58423 to 9.83566, sav
ing model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 1s - loss: 9.9614 - acc:
 0.2954 - val_loss: 9.8357 - val_acc: 0.3114
Epoch 3/20
6460/6680 [============================>.] - ETA: 0s - loss: 9.5477 - a
cc: 0.3605Epoch 00002: val_loss improved from 9.83566 to 9.66077, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 1s - loss: 9.5061 - acc:
 0.3630 - val_loss: 9.6608 - val_acc: 0.3329
Epoch 4/20
6660/6680 [============================>.] - ETA: 0s - loss: 9.3451 - a
cc: 0.3847Epoch 00003: val_loss did not improve
6680/6680 [==============================] - 1s - loss: 9.3509 - acc:
 0.3844 - val_loss: 9.6772 - val_acc: 0.3401
Epoch 5/20
6660/6680 [============================>.] - ETA: 0s - loss: 9.2615 - a
cc: 0.3986Epoch 00004: val_loss improved from 9.66077 to 9.62072, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 1s - loss: 9.2633 - acc:
 0.3985 - val_loss: 9.6207 - val_acc: 0.3485
Epoch 6/20
6460/6680 [============================>.] - ETA: 0s - loss: 9.1715 - a
cc: 0.4108Epoch 00005: val_loss did not improve
6680/6680 [==============================] - 1s - loss: 9.1853 - acc:
 0.4102 - val_loss: 9.6318 - val_acc: 0.3425
Epoch 7/20
6640/6680 [============================>.] - ETA: 0s - loss: 9.1106 - a
cc: 0.4173Epoch 00006: val_loss improved from 9.62072 to 9.45565, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 1s - loss: 9.1135 - acc:
 0.4169 - val_loss: 9.4557 - val_acc: 0.3581
Epoch 8/20
6660/6680 [============================>.] - ETA: 0s - loss: 8.9473 - a
cc: 0.4242Epoch 00007: val_loss improved from 9.45565 to 9.33405, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 1s - loss: 8.9454 - acc:
 0.4243 - val_loss: 9.3340 - val_acc: 0.3641
Epoch 9/20
6620/6680 [============================>.] - ETA: 0s - loss: 8.7640 - a
cc: 0.4381Epoch 00008: val_loss improved from 9.33405 to 9.13352, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] - 1s - loss: 8.7683 - acc:
 0.4377 - val_loss: 9.1335 - val_acc: 0.3760
Epoch 10/20
6480/6680 [============================>.] - ETA: 0s - loss: 8.6122 - a
cc: 0.4489Epoch 00009: val_loss improved from 9.13352 to 9.12358, savin
g model to saved_models/weights.best.VGG16.hdf5
```

```
6680/6680 [==============================] – 1s – loss: 8.6178 – acc:
 0.4491 – val_loss: 9.1236 – val_acc: 0.3760
Epoch 11/20
6480/6680 [=============================>.] – ETA: 0s – loss: 8.3715 – a
cc: 0.4616Epoch 00010: val_loss improved from 9.12358 to 8.96771, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] – 1s – loss: 8.3800 – acc:
 0.4608 – val_loss: 8.9677 – val_acc: 0.3892
Epoch 12/20
6540/6680 [=============================>.] – ETA: 0s – loss: 8.1764 – a
cc: 0.4769Epoch 00011: val_loss improved from 8.96771 to 8.74593, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] – 1s – loss: 8.1866 – acc:
 0.4762 – val_loss: 8.7459 – val_acc: 0.3832
Epoch 13/20
6540/6680 [=============================>.] – ETA: 0s – loss: 8.0025 – a
cc: 0.4885Epoch 00012: val_loss improved from 8.74593 to 8.68308, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] – 1s – loss: 8.0152 – acc:
 0.4879 – val_loss: 8.6831 – val_acc: 0.3868
Epoch 14/20
6580/6680 [=============================>.] – ETA: 0s – loss: 7.9681 – a
cc: 0.4957Epoch 00013: val_loss improved from 8.68308 to 8.66767, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] – 1s – loss: 7.9724 – acc:
 0.4954 – val_loss: 8.6677 – val_acc: 0.4084
Epoch 15/20
6580/6680 [=============================>.] – ETA: 0s – loss: 7.9515 – a
cc: 0.5015Epoch 00014: val_loss improved from 8.66767 to 8.58656, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] – 1s – loss: 7.9536 – acc:
 0.5012 – val_loss: 8.5866 – val_acc: 0.4084
Epoch 16/20
6620/6680 [=============================>.] – ETA: 0s – loss: 7.8652 – a
cc: 0.5063Epoch 00015: val_loss improved from 8.58656 to 8.51831, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] – 1s – loss: 7.8581 – acc:
 0.5067 – val_loss: 8.5183 – val_acc: 0.4048
Epoch 17/20
6540/6680 [=============================>.] – ETA: 0s – loss: 7.8358 – a
cc: 0.5090Epoch 00016: val_loss improved from 8.51831 to 8.51788, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] – 1s – loss: 7.8160 – acc:
 0.5102 – val_loss: 8.5179 – val_acc: 0.4084
Epoch 18/20
6660/6680 [=============================>.] – ETA: 0s – loss: 7.7490 – a
cc: 0.5122Epoch 00017: val_loss improved from 8.51788 to 8.40950, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] – 1s – loss: 7.7503 – acc:
 0.5120 – val_loss: 8.4095 – val_acc: 0.4204
Epoch 19/20
6660/6680 [=============================>.] – ETA: 0s – loss: 7.6504 – a
cc: 0.5135Epoch 00018: val_loss improved from 8.40950 to 8.36195, savin
g model to saved_models/weights.best.VGG16.hdf5
6680/6680 [==============================] – 1s – loss: 7.6477 – acc:
 0.5136 – val_loss: 8.3619 – val_acc: 0.4144
Epoch 20/20
```

```
         6640/6680 [============================>.] – ETA: 0s – loss: 7.5009 – a
         cc: 0.5212Epoch 00019: val_loss improved from 8.36195 to 8.20465, savin
         g model to saved_models/weights.best.VGG16.hdf5
         6680/6680 [==============================] – 1s – loss: 7.5085 – acc:
          0.5205 – val_loss: 8.2047 – val_acc: 0.4251
```

```
Out[21]: <keras.callbacks.History at 0x7efed8ce8cc0>
```

Load the Model with the Best Validation Loss

```
In [22]:  VGG16_model.load_weights('saved_models/weights.best.VGG16.hdf5')
```

Test the Model

Now, we can use the CNN to test how well it identifies breed within our test dataset of dog images. We print the test accuracy below.

```
In [23]:  # get index of predicted dog breed for each image in test set
          VGG16_predictions = [np.argmax(VGG16_model.predict(np.expand_dims(featur
          e, axis=0))) for feature in test_VGG16]

          # report test accuracy
          test_accuracy = 100*np.sum(np.array(VGG16_predictions)==np.argmax(test_t
          argets, axis=1))/len(VGG16_predictions)
          print('Test accuracy: %.4f%%' % test_accuracy)
```

```
          Test accuracy: 44.7368%
```

Predict Dog Breed with the Model

```
In [21]:  from extract_bottleneck_features import *

          def VGG16_predict_breed(img_path):
              # extract bottleneck features
              bottleneck_feature = extract_VGG16(path_to_tensor(img_path))
              # obtain predicted vector
              predicted_vector = VGG16_model.predict(bottleneck_feature)
              # return dog breed that is predicted by the model
              return dog_names[np.argmax(predicted_vector)]
```

Step 5: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

In Step 4, we used transfer learning to create a CNN using VGG-16 bottleneck features. In this section, you must use the bottleneck features from a different pre-trained model. To make things easier for you, we have pre-computed the features for all of the networks that are currently available in Keras:

- VGG-19 (https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogVGG19Data.npz) bottleneck features
- ResNet-50 (https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogResnet50Data.npz) bottleneck features
- Inception (https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogInceptionV3Data.npz) bottleneck features
- Xception (https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogXceptionData.npz) bottleneck features

The files are encoded as such:

```
Dog{network}Data.npz
```

where `{network}`, in the above filename, can be one of `VGG19`, `Resnet50`, `InceptionV3`, or `Xception`. Pick one of the above architectures, download the corresponding bottleneck features, and store the downloaded file in the `bottleneck_features/` folder in the repository.

(IMPLEMENTATION) Obtain Bottleneck Features

In the code block below, extract the bottleneck features corresponding to the train, test, and validation sets by running the following:

```
bottleneck_features = np.load('bottleneck_features/Dog{network}Data.npz')
train_{network} = bottleneck_features['train']
valid_{network} = bottleneck_features['valid']
test_{network} = bottleneck_features['test']
```

```
In [14]:   ### TODO: Obtain bottleneck features from another pre-trained CNN.
           bottleneck_features = np.load('bottleneck_features/DogXceptionData.npz')
           train_xception = bottleneck_features['train']
           valid_xception = bottleneck_features['valid']
           test_xception = bottleneck_features['test']
```

(IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
<your model's name>.summary()
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer:

I decided to use the Xception model to take advantage of transfer learning. The model architecture takes the feature extractions from Xception, then it adds a global average pooling layer from the Xception output and finally it has an output layer to predict the dog breeds with a softmax activation function.

This model seemed reasonable for this task given that the Xception algorithm has had a lot of success in image classification tasks with state of the art performance. It also seemed like a reasonable start given the above example that used VGG16.

```
In [15]:  ### TODO: Define your architecture.
          xception_model = Sequential()

          xception_model.add(GlobalAveragePooling2D(input_shape=train_xception.sha
          pe[1:]))
          xception_model.add(Dense(133, activation='softmax'))

          xception_model.summary()
```

```
_____
Layer (type)                     Output Shape              Param #
=================================================================
global_average_pooling2d_1 (  (None, 2048)              0
_____
dense_2 (Dense)                  (None, 133)               272517
=================================================================
Total params: 272,517
Trainable params: 272,517
Non-trainable params: 0
_____
```

(IMPLEMENTATION) Compile the Model

```
In [27]:  ### TODO: Compile the model.
          xception_model.compile(loss='categorical_crossentropy', optimizer='rmspr
          op', metrics=['accuracy'])
```

(IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to augment the training data (https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.

```
In [28]:  ### TODO: Train the model.
          checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.xcept
          ion.hdf5',
                                         verbose=1, save_best_only=True)

          xception_model.fit(train_xception, train_targets,
                  validation_data=(valid_xception, valid_targets),
                  epochs=20, batch_size=20, callbacks=[checkpointer], verbose=1)
```

```
Train on 6680 samples, validate on 835 samples
Epoch 1/20
6660/6680 [==============================>.] - ETA: 0s - loss: 1.0535 - a
cc: 0.7324Epoch 00000: val_loss improved from inf to 0.52580, saving mo
del to saved_models/weights.best.xception.hdf5
6680/6680 [==============================] - 3s - loss: 1.0527 - acc:
 0.7323 - val_loss: 0.5258 - val_acc: 0.8216
Epoch 2/20
6580/6680 [==============================>.] - ETA: 0s - loss: 0.4029 - a
cc: 0.8722Epoch 00001: val_loss improved from 0.52580 to 0.47015, savin
g model to saved_models/weights.best.xception.hdf5
6680/6680 [==============================] - 3s - loss: 0.4024 - acc:
 0.8725 - val_loss: 0.4702 - val_acc: 0.8539
Epoch 3/20
6620/6680 [==============================>.] - ETA: 0s - loss: 0.3217 - a
cc: 0.8977Epoch 00002: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.3207 - acc:
 0.8982 - val_loss: 0.4794 - val_acc: 0.8503
Epoch 4/20
6540/6680 [==============================>.] - ETA: 0s - loss: 0.2747 - a
cc: 0.9122Epoch 00003: val_loss improved from 0.47015 to 0.45458, savin
g model to saved_models/weights.best.xception.hdf5
6680/6680 [==============================] - 2s - loss: 0.2770 - acc:
 0.9127 - val_loss: 0.4546 - val_acc: 0.8611
Epoch 5/20
6600/6680 [==============================>.] - ETA: 0s - loss: 0.2483 - a
cc: 0.9229Epoch 00004: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.2472 - acc:
 0.9237 - val_loss: 0.5084 - val_acc: 0.8479
Epoch 6/20
6620/6680 [==============================>.] - ETA: 0s - loss: 0.2208 - a
cc: 0.9334Epoch 00005: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.2206 - acc:
 0.9335 - val_loss: 0.5357 - val_acc: 0.8563
Epoch 7/20
6600/6680 [==============================>.] - ETA: 0s - loss: 0.1963 - a
cc: 0.9382Epoch 00006: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.1963 - acc:
 0.9380 - val_loss: 0.5573 - val_acc: 0.8551
Epoch 8/20
6580/6680 [==============================>.] - ETA: 0s - loss: 0.1810 - a
cc: 0.9422Epoch 00007: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.1809 - acc:
 0.9425 - val_loss: 0.5368 - val_acc: 0.8659
Epoch 9/20
6560/6680 [==============================>.] - ETA: 0s - loss: 0.1630 - a
cc: 0.9497Epoch 00008: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.1631 - acc:
 0.9499 - val_loss: 0.5581 - val_acc: 0.8599
Epoch 10/20
6560/6680 [==============================>.] - ETA: 0s - loss: 0.1516 - a
cc: 0.9524Epoch 00009: val_loss did not improve
6680/6680 [==============================] - 2s - loss: 0.1500 - acc:
 0.9530 - val_loss: 0.5563 - val_acc: 0.8539
Epoch 11/20
6600/6680 [==============================>.] - ETA: 0s - loss: 0.1392 - a
cc: 0.9577Epoch 00010: val_loss did not improve
```

```
6680/6680 [==============================] - 3s - loss: 0.1380 - acc:
  0.9582 - val_loss: 0.5741 - val_acc: 0.8587
Epoch 12/20
6620/6680 [=============================>.] - ETA: 0s - loss: 0.1281 - a
cc: 0.9609Epoch 00011: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.1284 - acc:
  0.9611 - val_loss: 0.5956 - val_acc: 0.8635
Epoch 13/20
6580/6680 [=============================>.] - ETA: 0s - loss: 0.1209 - a
cc: 0.9644Epoch 00012: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.1200 - acc:
  0.9647 - val_loss: 0.5901 - val_acc: 0.8671
Epoch 14/20
6660/6680 [=============================>.] - ETA: 0s - loss: 0.1111 - a
cc: 0.9676Epoch 00013: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.1111 - acc:
  0.9675 - val_loss: 0.6239 - val_acc: 0.8515
Epoch 15/20
6620/6680 [=============================>.] - ETA: 0s - loss: 0.1034 - a
cc: 0.9710Epoch 00014: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.1038 - acc:
  0.9708 - val_loss: 0.6300 - val_acc: 0.8551
Epoch 16/20
6540/6680 [=============================>.] - ETA: 0s - loss: 0.0954 - a
cc: 0.9732Epoch 00015: val_loss did not improve
6680/6680 [==============================] - 3s - loss: 0.0962 - acc:
  0.9731 - val_loss: 0.6645 - val_acc: 0.8587
Epoch 17/20
6560/6680 [=============================>.] - ETA: 0s - loss: 0.0923 - a
cc: 0.9736Epoch 00016: val_loss did not improve
6680/6680 [==============================] - 2s - loss: 0.0929 - acc:
  0.9732 - val_loss: 0.6580 - val_acc: 0.8599
Epoch 18/20
6560/6680 [=============================>.] - ETA: 0s - loss: 0.0876 - a
cc: 0.9753Epoch 00017: val_loss did not improve
6680/6680 [==============================] - 2s - loss: 0.0873 - acc:
  0.9754 - val_loss: 0.6934 - val_acc: 0.8575
Epoch 19/20
6620/6680 [=============================>.] - ETA: 0s - loss: 0.0823 - a
cc: 0.9770Epoch 00018: val_loss did not improve
6680/6680 [==============================] - 2s - loss: 0.0818 - acc:
  0.9771 - val_loss: 0.6968 - val_acc: 0.8515
Epoch 20/20
6660/6680 [=============================>.] - ETA: 0s - loss: 0.0772 - a
cc: 0.9787Epoch 00019: val_loss did not improve
6680/6680 [==============================] - 2s - loss: 0.0770 - acc:
  0.9787 - val_loss: 0.6914 - val_acc: 0.8563
```

Out[28]: `<keras.callbacks.History at 0x7efed8af1f60>`

(IMPLEMENTATION) Load the Model with the Best Validation Loss

In [16]:
```
### TODO: Load the model weights with the best validation loss.
xception_model.load_weights('saved_models/weights.best.xception.hdf5')
```

(IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 60%.

```
In [17]:  ### TODO: Calculate classification accuracy on the test dataset.
          # get index of predicted dog breed for each image in test set
          xception_predictions =
          [np.argmax(xception_model.predict(np.expand_dims(feature, axis=0))) for
          feature in test_xception]

          # report test accuracy
          test_accuracy = 100*np.sum(np.array(xception_predictions)==np.argmax(tes
          t_targets, axis=1))/len(xception_predictions)
          print('Test accuracy: %.4f%%' % test_accuracy)
```

```
Test accuracy: 85.2871%
```

(IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (`Affenpinscher`, `Afghan_hound`, etc) that is predicted by your model.

Similar to the analogous function in Step 5, your function should have three steps:

1. Extract the bottleneck features corresponding to the chosen CNN model.
2. Supply the bottleneck features as input to the model to return the predicted vector. Note that the argmax of this prediction vector gives the index of the predicted dog breed.
3. Use the `dog_names` array defined in Step 0 of this notebook to return the corresponding breed.

The functions to extract the bottleneck features can be found in `extract_bottleneck_features.py`, and they have been imported in an earlier code cell. To obtain the bottleneck features corresponding to your chosen CNN architecture, you need to use the function

```
extract_{network}
```

where `{network}`, in the above filename, should be one of `VGG19`, `Resnet50`, `InceptionV3`, or `Xception`.

```
In [18]:  ### TODO: Write a function that takes a path to an image as input
          ### and returns the dog breed that is predicted by the model.
          def xception_predict_breed(img_path):
              # extract bottleneck features
              bottleneck_feature = extract_Xception(path_to_tensor(img_path))
              # obtain predicted vector
              predicted_vector = xception_model.predict(bottleneck_feature)
              # return dog breed that is predicted by the model
              return dog_names[np.argmax(predicted_vector)]
```

Step 6: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then,

- if a **dog** is detected in the image, return the predicted breed.
- if a **human** is detected in the image, return the resembling dog breed.
- if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `dog_detector` functions developed above. You are **required** to use your CNN from Step 5 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!


```
hello, human!
```

```
You look like a ...
Chinese_shar-pei
```

(IMPLEMENTATION) Write your Algorithm

```
In [42]:   ### TODO: Write your algorithm.
           ### Feel free to use as many code cells as needed.
           def predict_breed(img_path):

               is_person = face_detector(img_path)
               is_dog = dog_detector(img_path)

               if is_person and not is_dog:
                   breed = xception_predict_breed(img_path)
                   return "Hi! You look like a {}".format(breed)
               elif is_dog:
                   breed = xception_predict_breed(img_path)
                   return "This dog is a {}".format(breed)
               else:
                   return "We could not detect a person or a dog in the image. Plea
           se try again with another image."
```

Step 7: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that **you** look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

(IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer:

In [43]:
```
## TODO: Execute your algorithm from Step 6 on
## at least 6 images on your computer.
## Feel free to use as many code cells as needed.

img_path = 'images/person_1.JPG'
img = cv2.imread( img_path )
plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
result = predict_breed(img_path)
print( result )
```

Hi! You look like a Chesapeake_bay_retriever

In [45]:
```
img_path = 'images/person_caro.jpg'
img = cv2.imread( img_path )
plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
result = predict_breed(img_path)
print( result )
```

Hi! You look like a Dachshund

In [49]:
```
img_path = 'images/person_sebas.jpg'
img = cv2.imread( img_path )
plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
result = predict_breed(img_path)
print( result )
```

Hi! You look like a Dachshund

In [46]:
```
img_path = 'images/dog_emma.JPG'
img = cv2.imread( img_path )
plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
result = predict_breed(img_path)
print( result )
```

This dog is a Parson_russell_terrier

```
In [48]:   img_path = 'images/dog_aisha.jpg'
           img = cv2.imread( img_path )
           plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
           result = predict_breed(img_path)
           print( result )
```

This dog is a Boston_terrier


```
In [50]:   img_path = 'images/dog_aggie.JPG'
           img = cv2.imread( img_path )
           plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
           result = predict_breed(img_path)
           print( result )
```

This dog is a Labrador_retriever

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

The output for classifying dogs is actually pretty good. It managed to classify my 3 dogs correctly. The algorithm did not perform well for human faces though.

There is a lot of room for improving the performance and usability of the algorithm. These are some of the things to improve:

- If the picture contains more than 1 human face, crop each face and classify each person independently
- Improve the human detection algorithm with a better model using deep learning.
- The human and dog detection algorithms could return a probabily value instead of boolean, so as to let the algorithm choose a measure of uncertainty.
- Improve the dog classification algorithm by increasing training time

```
In [ ]:
```